



The City College
of New York

CSC 36000: Modern Distributed Computing *with AI Agents*

By Saptarashmi Bandyopadhyay

Email: sbandyopadhyay@ccny.cuny.edu

Assistant Professor of Computer Science

City College of New York and Graduate Center at City University of New York

November 24, 2025 CSC 36000

Today's Lecture

Data Distribution Algorithms

- Anti-Entropy
- Gossip Protocol
- Coding Demonstration

MapReduce

- Primitive Operations: Map, Shuffle, Reduce
- Coding Demonstration

Data in Distributed Computing

Propagating Data in Distributed Systems

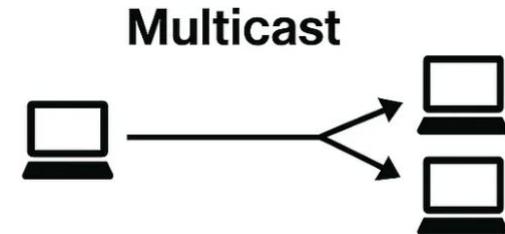
How do we reliably update data across 10,000+ nodes in a decentralized network?

Traditional Reliable Multicast relies on acknowledgments (**ACKs**) or negative acknowledgments (**NAKs**).

As the number of receivers **N** grows, the sender becomes swamped with feedback messages.

Centralized directories (like basic DNS) create bottlenecks and single points of failure.

We need a mechanism that relies only on *local* information yet achieves *global* convergence!



Data Distribution as an Infectious Disease

Epidemic Theory:

- Modeled after the spread of infectious diseases.
- Instead of a cold, we are spreading updates or information.

Goal: "Infect" 100% of the population with the update as fast as possible.

Key Terminology:

- **Infected:** A node that holds the update and is willing to share it.
- **Susceptible:** A node that has not yet seen the update.
- **Removed:** A node that has the update but is no longer sharing it (lazy/inactive).



Anti-Entropy

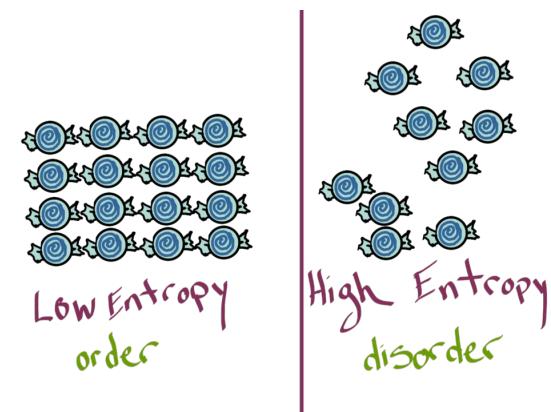
Nodes regularly choose a random neighbor to exchange state with.

Goal: Resolving inconsistencies (entropy) between nodes.

Three Approaches:

- **Push:** P sends updates to Q. (Efficient when few are infected).
- **Pull:** P asks Q for updates. (Efficient when many are infected).
- **Push-Pull:** Both exchange updates. (Optimal convergence: $O(\log N)$ rounds).

Anti-Entropy Guarantees eventual consistency but can be bandwidth intensive (sending large state differences).



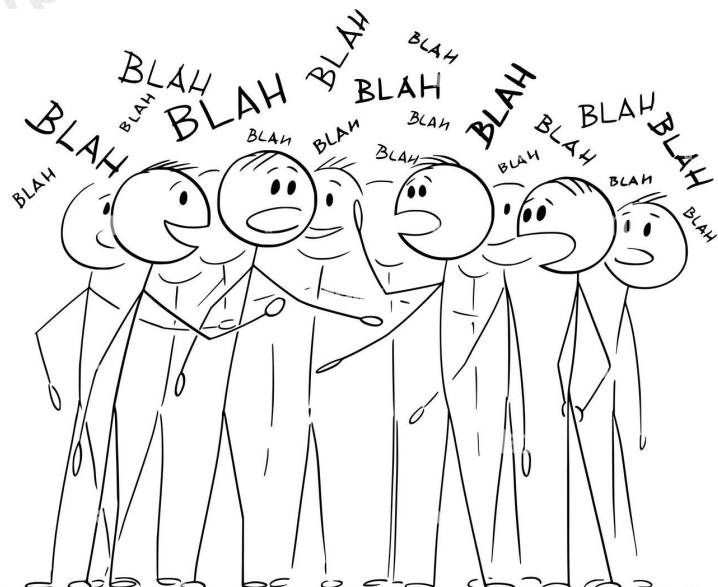
Gossip Method Protocol

If a node receives a new update, it tries to push it to a random neighbor.

If the neighbor already knows the update, the sender loses interest. The sender stops spreading the rumor with probability $1 / k$.

Pros: Extremely rapid distribution; low bandwidth when the system is dormant.

Cons: Probabilistic. There is a small chance some nodes remain Susceptible (never receive the update) if the "gossip dies out" too quickly



Deletion in Distributed Computing

Why is deleting hard?

- In a stateless system, a "deletion" looks like a missing file.
- If Node A deletes Item X, and later syncs with Node B (who still has X), Node B might "update" Node A by sending X back.



The Solution: Death Certificates

- Don't delete the item immediately.
- Replace the item with a Death Certificate: A timestamped record stating "Item X is deleted."
- Spread the certificate via gossip like any other update.

Certificates are eventually cleaned up after a maximum propagation time.

Hands-on Gossip Deletion Example

[https://drive.google.com/file/d/14piiD1Pem_o5es4b9Pc7QU
SI40iWmLa5/view?usp=sharing](https://drive.google.com/file/d/14piiD1Pem_o5es4b9Pc7QU SI40iWmLa5/view?usp=sharing)

MapReduce

Saptarashmi Bandyopadhyay

Working with Big Data

In modern Distributed Computing and AI, we're often working with internet-scale datasets. How do we process so much data?

Traditional Approach (Serial Processing):

- Store data on one massive storage server (SAN/NAS).
- Pull data across the network to a supercomputer.
- Process line-by-line.

The Bottlenecks:

- Network Bandwidth: Moving 10PB takes forever.
- Single Point of Failure: If the processor crashes at 99%, you start over.
- Disk I/O: Reading linearly is slow.



Divide and Conquer

Google addressed this problem back in 2004:

- Don't buy one supercomputer. *Buy 1,000 cheap commodity servers.*
- Data Locality: Don't move data to the CPU. *Move the code to the data.*

The Abstraction:

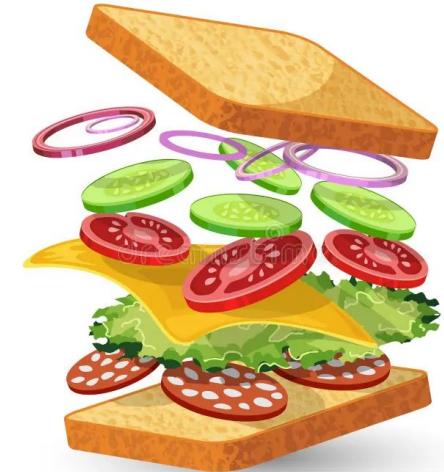
- Hide the messy details (network failures, disk crashes, load balancing).
- Expose two simple functions to the programmer: Map and Reduce.

Analogy: Sandwich Shop

Goal: Make 1,000 Sandwiches as fast as possible.

Phase 1: MAP (The Prep Cooks)

- 5 cooks working in parallel.
- They don't make full sandwiches. They just output intermediate parts.
- Output: Piles of "Sliced Tomato," "Chopped Lettuce," "Cooked Bacon."



Phase 2: SHUFFLE (The Runners)

- The messy part. Runners grab all tomatoes from all 5 cooks and put them in the "Tomato Bin."
- Grouping: Ensure all instances of the same ingredient end up in the same place.

Phase 3: REDUCE (The Assemblers)

- One chef stands at the Tomato bin, one at the Bacon bin.
- They process the piles into the final result.

The Technical Approach

Input: Split into chunks (usually 64MB or 128MB).

MAP: (Key, Value) -> List(Key, Value)

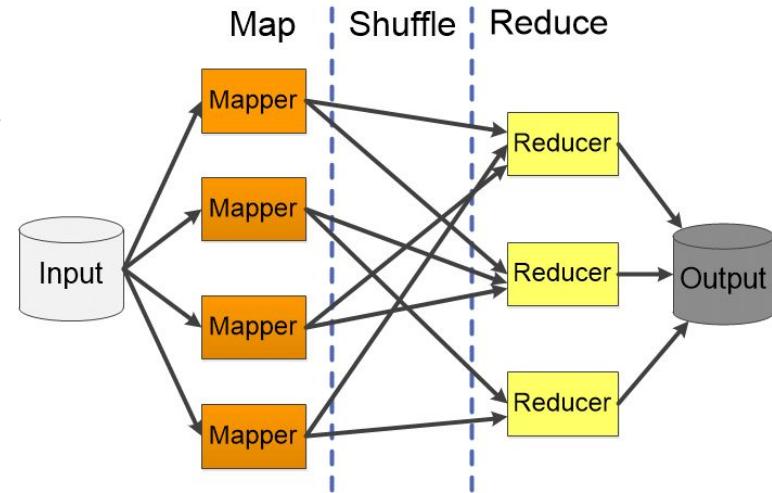
- Takes raw data, filters/parses it, outputs intermediate pairs.
- Example: (Doc1, "Hello World") -> [("Hello", 1), ("World", 1)]

SHUFFLE & SORT: (Key, Value) -> (Key, List[Values])

- The system groups all values associated with the same key.
- Example: ("Hello", [1, 1, 1, 1])

REDUCE: (Key, List[Values]) -> (Key, Result)

- Aggregates the list into a single output.
- Example: ("Hello", 4)



Example: Word Count

Input Data: "Deer Bear River" "Car Car River" "Deer Car Bear"

Map Phase (Parallel):

- Worker 1: (Deer, 1), (Bear, 1), (River, 1)
- Worker 2: (Car, 1), (Car, 1), (River, 1)
- Worker 3: (Deer, 1), (Car, 1), (Bear, 1)

Shuffle Phase (Network Intensive): Moves data so keys group together.

- Bear: [1, 1]
- Car: [1, 1, 1]
- Deer: [1, 1]
- River: [1, 1]

Reduce Phase: Sum the lists: Bear: 2, Car: 3, Deer: 2, River: 2

Overhead

Synchronization Barrier

- The Reduce phase cannot start until the slowest Map worker is finished.
- Analogy: You can't count total "Tomatoes" until the last prep cook finishes chopping.

Network Congestion:

- The "Shuffle" phase requires moving massive amounts of data between servers simultaneously.

Fault Tolerance:

- What if Worker 5 dies? The Master node must detect it and re-assign that chunk of data to Worker 2.

Hands-On MapReduce Demonstration

[https://drive.google.com/file/d/1PpZDwDgH_0-R4ArOOFIUdybn82GQUm2 /view?usp=sharing](https://drive.google.com/file/d/1PpZDwDgH_0-R4ArOOFIUdybn82GQUm2/view?usp=sharing)

Questions?

Saptarashmi Bandyopadhyay